

# Sentiment Analysis

SDS 322E  
October 17, 2025

Dr. Lydia R. Lucchesi  
Department of Statistics and Data Sciences  
The University of Texas at Austin

# Week 8

	Monday	Tuesday	Wednesday	Thursday	Friday
Lecture	Project 1 Workday		Introduction to Tidy Text		Sentiment Analysis
Other	Lab 6	Office Hours (3-5PM)	Office Hours (3-5PM)	Project 1 Due	Pre-Lab 7 Quiz Due

In text analysis, what do we call a meaningful unit of text (most often a word)?

Which function from the **tidytext** R package do we use to tokenize text data?

- a. `tokenize_text()`
- b. `split_text()`
- c. `unnest_tokens()`

Which of the three lexicons in the **stop\_words** tibble contains the most stop words?

- a. onix
- b. SMART
- c. snowball

**Positive or negative sentiment?**

When the plane landed, I could see from my little window that the weather was gloomy. I trudged over to baggage claim, where I waited for many minutes, before it became frustratingly clear that my luggage had been lost.

**Positive or negative sentiment?**

The landing was smooth, and when the plane parked at the gate, I could barely contain my excitement. My favorite place in the world was beckoning to me.

# The **sentiments** dataset in the **tidytext** R package

```
> sentiments |>
+   slice(2677:2687)
# A tibble: 11 × 2
   word      sentiment
  <chr>    <chr>
1 gloomy   negative
2 glorify  positive
3 glorious positive
4 gloriously positive
5 glory    positive
6 glow     positive
7 glower   negative
8 glowing  positive
9 glowingly positive
10 glum     negative
11 glut     negative
```

## Details

This lexicon was first published in:

Minqing Hu and Bing Liu, "Mining and summarizing customer reviews.", Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery & Data Mining (KDD-2004), Seattle, Washington, USA, Aug 22-25, 2004.

Words with non-ASCII characters were removed.



```

stories <- tibble(story = read_lines("plane_stories.txt"))

smart_stop_words <- stop_words |>
  filter(lexicon == "SMART")

stories |>
  filter(str_detect(story, "[a-zA-Z]+")) |>
  mutate(id = row_number()) |>
  relocate(id) |>
  unnest_tokens(word, story) |>
  anti_join(smart_stop_words, join_by(word))

```

```

# A tibble: 25 x 2
      id word
  <int> <chr>
1     1 plane
2     1 landed
3     1 window
4     1 weather
5     1 gloomy
6     1 trudged
7     1 baggage
8     1 claim
9     1 waited
10    1 minutes

```

The next step is to add on sentiment information. If we want to keep only words that exist in both **stories** and **sentiments**, which join function should we use?

```

stories <- tibble(story = read_lines("plane_stories.txt"))

smart_stop_words <- stop_words |>
  filter(lexicon == "SMART")

stories |>
  filter(str_detect(story, "[a-zA-Z]+")) |>
  mutate(id = row_number()) |>
  relocate(id) |>
  unnest_tokens(word, story) |>
  anti_join(smart_stop_words, join_by(word)) |>
  inner_join(sentiments, join_by(word), multiple = "all")

```

```

# A tibble: 8 x 3
      id word      sentiment
  <int> <chr>    <chr>
1     1 gloomy    negative
2     1 frustratingly negative
3     1 clear     positive
4     1 lost      negative
5     2 smooth    positive
6     2 excitement positive
7     2 favorite    positive
8     2 beckoning  positive

```



```

sentiments |>
  count(word) |>
  filter(n > 1)

```

```

# A tibble: 3 x 2
  word      n
  <chr>    <int>
1 envious    2
2 enviously  2
3 enviousness 2

```

# Analyzing sentiment in texts from Project Gutenberg

```
library(tidyverse)
library(tidytext)
# So we can download public domain works from Project Gutenberg
library(gutenbergr)

# Get SMART stop word lexicon
smart_stop_words <- stop_words |>
  filter(lexicon == "SMART")

# Download the book Frankenstein
book <- gutenbergr_download(84)

# Tidy and tokenize, remove stop words, and add sentiment
book |>
  mutate(chapter = cumsum(str_detect(text, "^Chapter")))) |>
  filter(chapter > 0) |>
  filter(!str_detect(text, "^Chapter")) |>
  filter(str_detect(text, "[a-zA-Z]+")) |>
  unnest_tokens(word, text) |>
  anti_join(smart_stop_words, join_by(word)) |>
  inner_join(sentiments, join_by(word), multiple = "all")
```

# Analyzing sentiment in texts from Project Gutenberg

```
book |>
  mutate(chapter = cumsum(str_detect(text, "^Chapter"))) |>
  filter(chapter > 0) |>
  filter(!str_detect(text, "^Chapter")) |>
  filter(str_detect(text, "[a-zA-Z]+")) |>
  unnest_tokens(word, text) |>
  anti_join(smart_stop_words, join_by(word)) |>
  inner_join(sentiments, join_by(word), multiple = "all") |>
  group_by(chapter, sentiment) |>
  summarize(n = n(), .groups = "drop") |>
  pivot_wider(names_from = "sentiment", values_from = "n") |>
  mutate(proportion = positive / (positive + negative))
```

# Analyzing sentiment in texts from Project Gutenberg


```
book |>
  mutate(chapter = cumsum(str_detect(text, "^Chapter"))) |>
  filter(chapter > 0) |>
  filter(!str_detect(text, "^Chapter")) |>
  filter(str_detect(text, "[a-zA-Z]+")) |>
  unnest_tokens(word, text) |>
  anti_join(smart_stop_words, join_by(word)) |>
  inner_join(sentiments, join_by(word), multiple = "all") |>
  group_by(chapter, sentiment) |>
  summarize(n = n(), .groups = "drop") |>
  pivot_wider(names_from = "sentiment", values_from = "n") |>
  mutate(proportion = positive / (positive + negative))
```

```
# A tibble: 24 × 4
```

	chapter	negative	positive	proportion
	<int>	<int>	<int>	<dbl>
1	1	60	94	0.610
2	2	84	112	0.571
3	3	88	88	0.5
4	4	104	115	0.525
5	5	109	76	0.411
6	6	80	126	0.612
7	7	214	115	0.350
8	8	227	91	0.286
9	9	139	92	0.398
10	10	153	70	0.314

# Analyzing sentiment in texts from Project Gutenberg

```
book |>
  mutate(chapter = cumsum(str_detect(text, "^Chapter"))) |>
  filter(chapter > 0) |>
  filter(!str_detect(text, "^Chapter")) |>
  filter(str_detect(text, "[a-zA-Z]+")) |>
  unnest_tokens(word, text) |>
  anti_join(smart_stop_words, join_by(word)) |>
  inner_join(sentiments, join_by(word), multiple = "all") |>
  group_by(chapter, sentiment) |>
  summarize(n = n(), .groups = "drop") |>
  pivot_wider(names_from = "sentiment", values_from = "n") |>
  mutate(proportion = positive / (positive + negative)) |>
  ggplot(aes(x = chapter, y = proportion)) +
    geom_hline(yintercept = .5, linewidth = .25, color = "gray") +
    geom_line() +
    geom_point() +
    scale_x_continuous(breaks = 1:24) +
    scale_y_continuous(limits = c(0.2, 0.8),
                      breaks = seq(0.2, 0.8, 0.1)) +
    theme_light() +
    theme(panel.grid.minor = element_blank(),
          panel.grid.major = element_blank()) +
    labs(x = "Chapter",
         y = "Proportion of Positive Sentiment Words")
```



# ggwordcloud R package

```
book |>
  mutate(chapter = cumsum(str_detect(text, "^Chapter"))) |>
  filter(chapter > 0) |>
  filter(!str_detect(text, "^Chapter")) |>
  filter(str_detect(text, "[a-zA-Z]+")) |>
  unnest_tokens(word, text) |>
  anti_join(smart_stop_words, join_by(word)) |>
  inner_join(sentiments, join_by(word), multiple = "all") |>
  filter(chapter %in% 1:5) |>
  group_by(word, sentiment) |>
  summarize(n = n(), .groups = "drop") |>
  arrange(desc(n)) |>
  slice(1:15) |>
  ggplot(aes(label = word, size = n, color = sentiment)) +
  geom_text_wordcloud()
```

